

COMPUTER SCIENCE

Database Management System

Transaction & Concurrency Control

Lecture_10



Vijay Agarwal sir

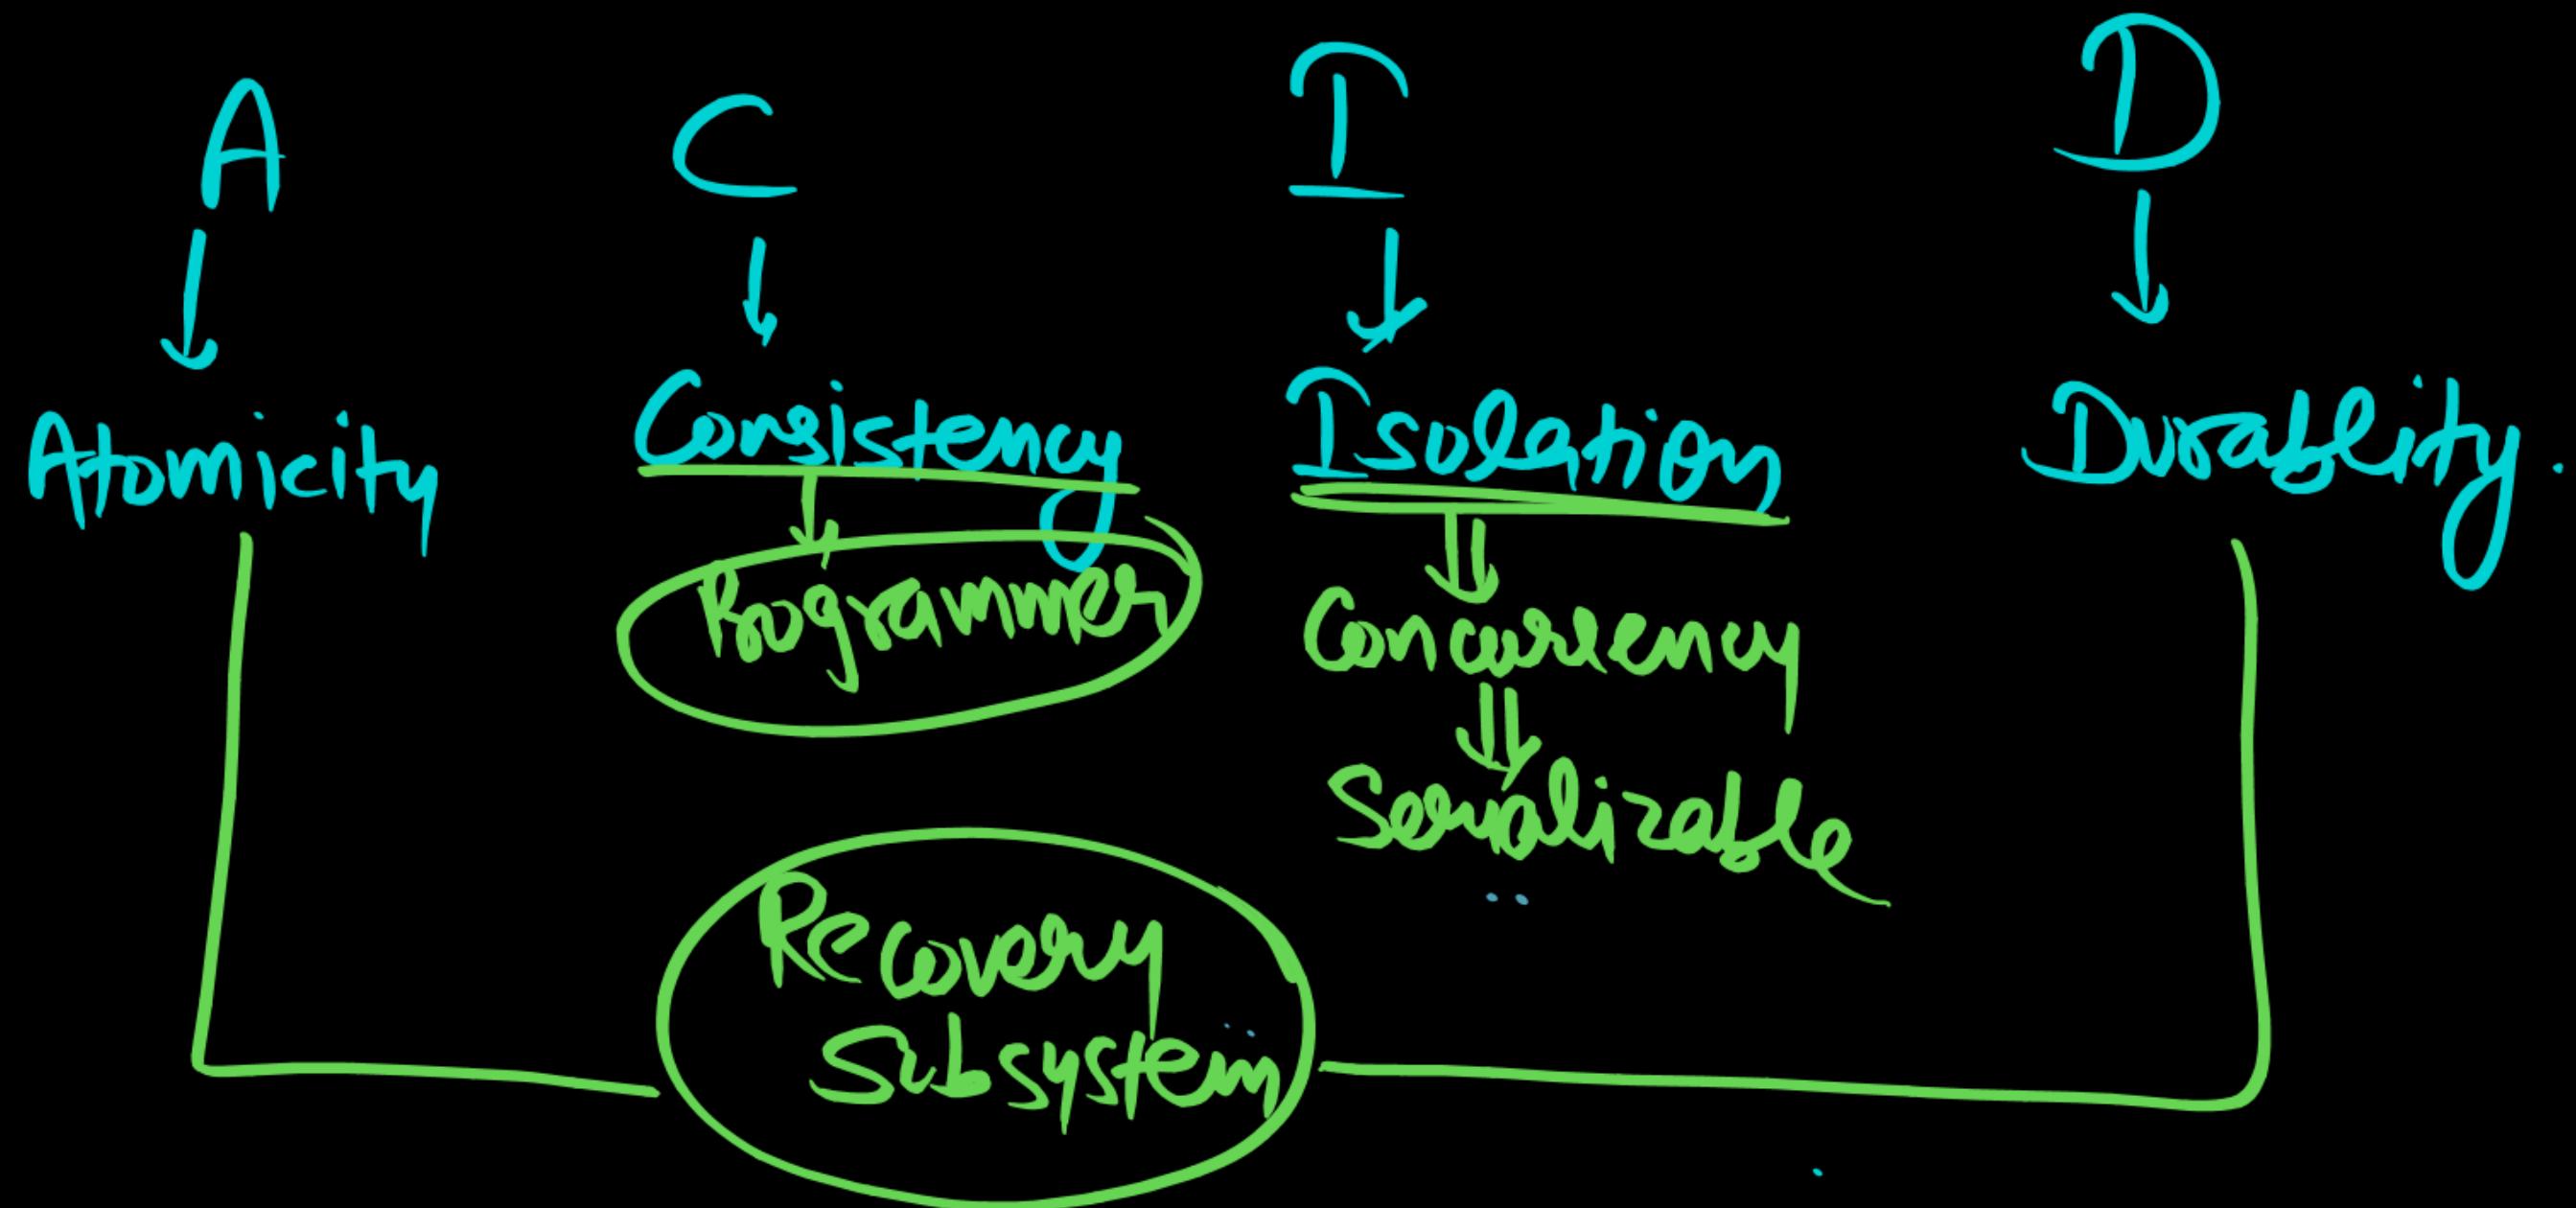
A graphic of a construction barrier made of orange and white striped panels, topped with two yellow bollards. It is positioned on the left side of the slide.

**TOPICS
TO BE
COVERED**

01

Lock Based Protocol

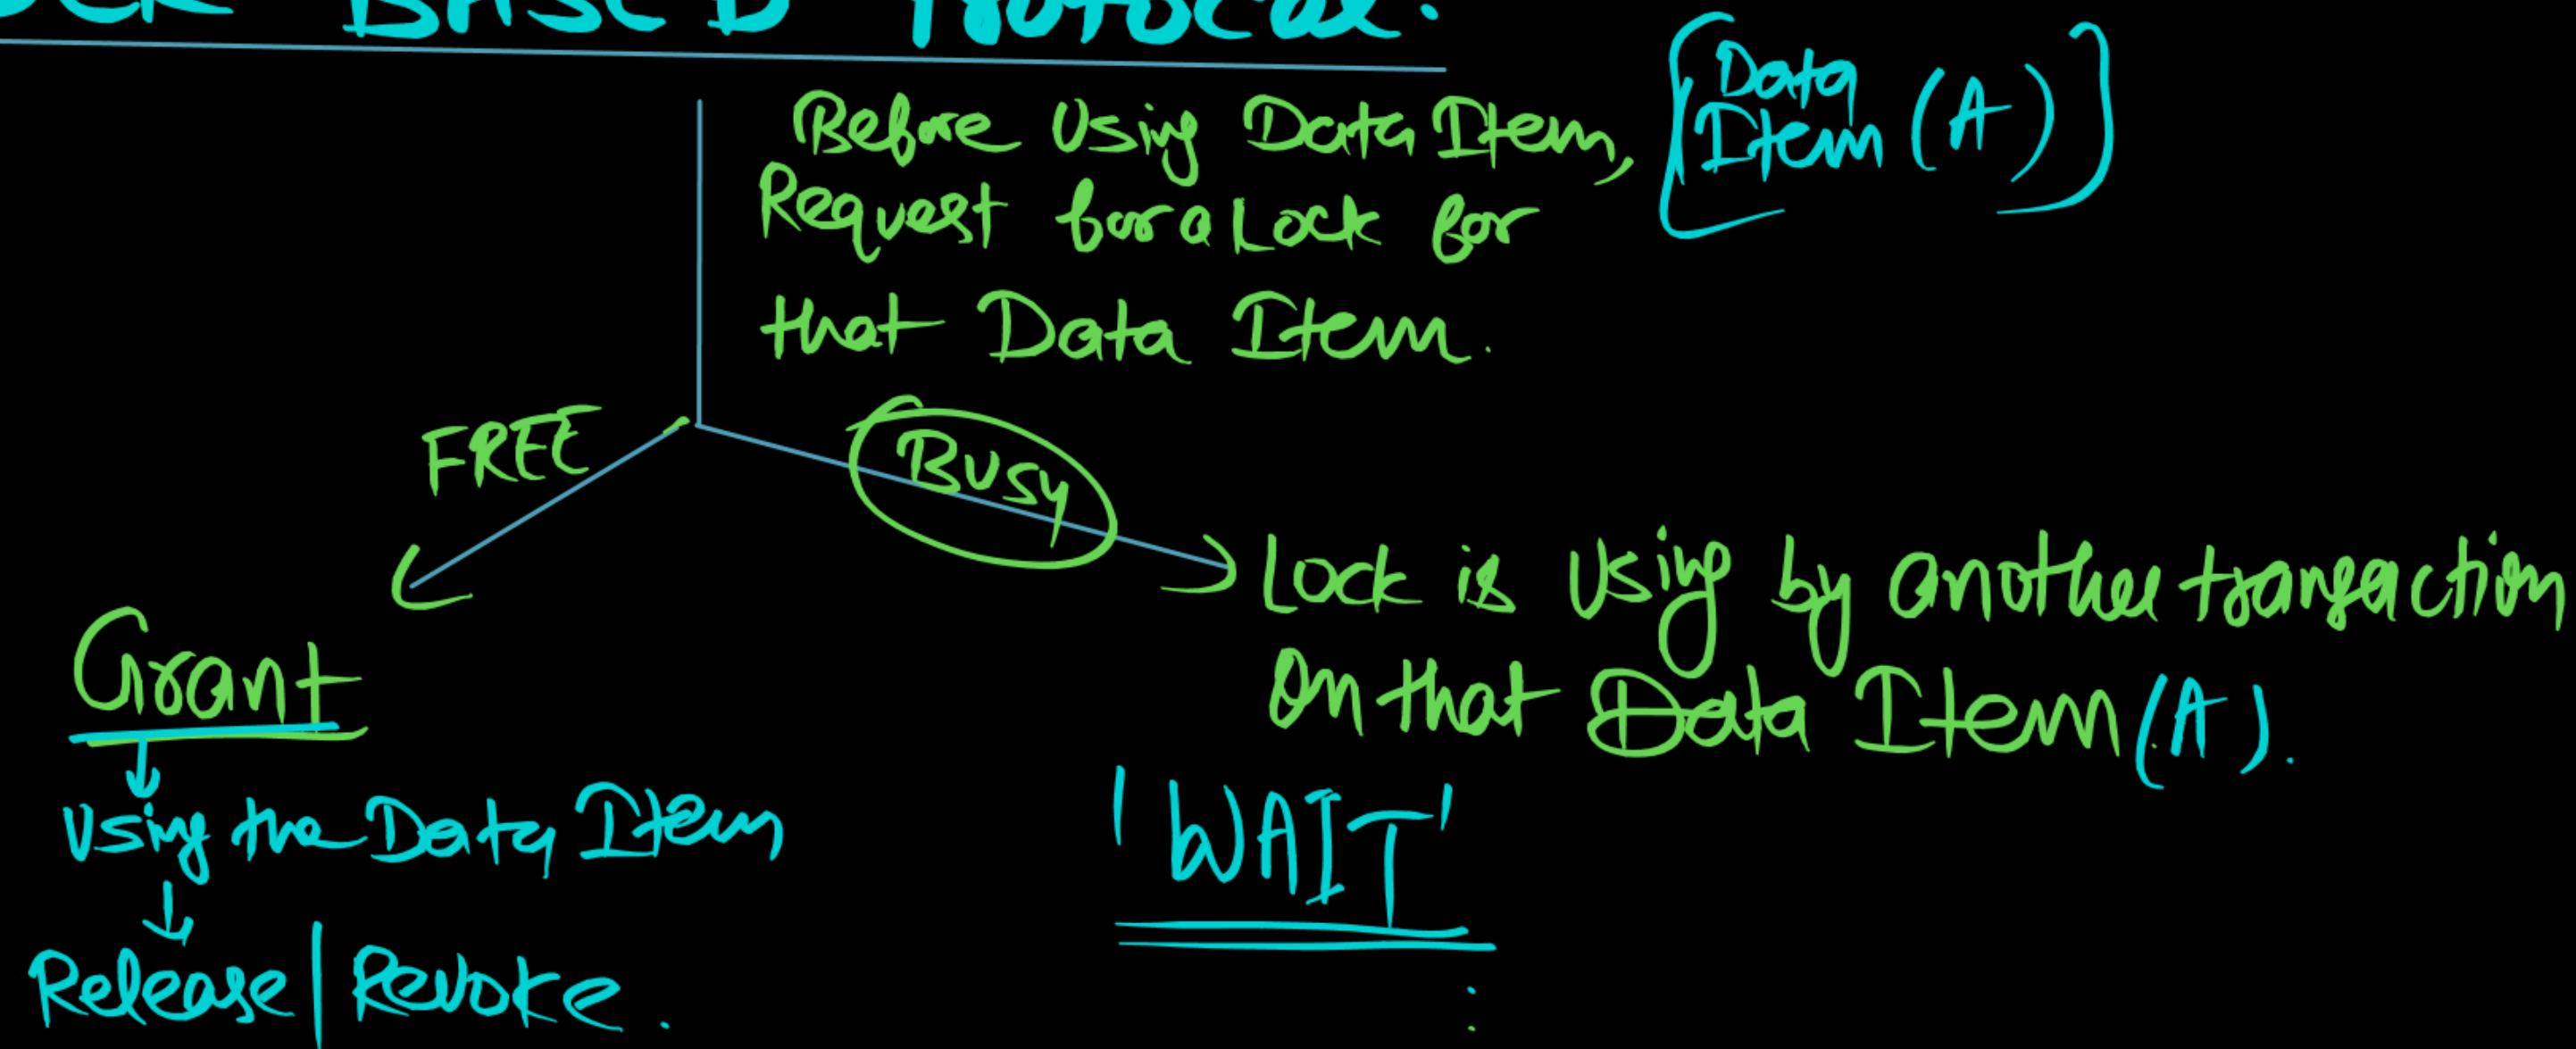
Transaction

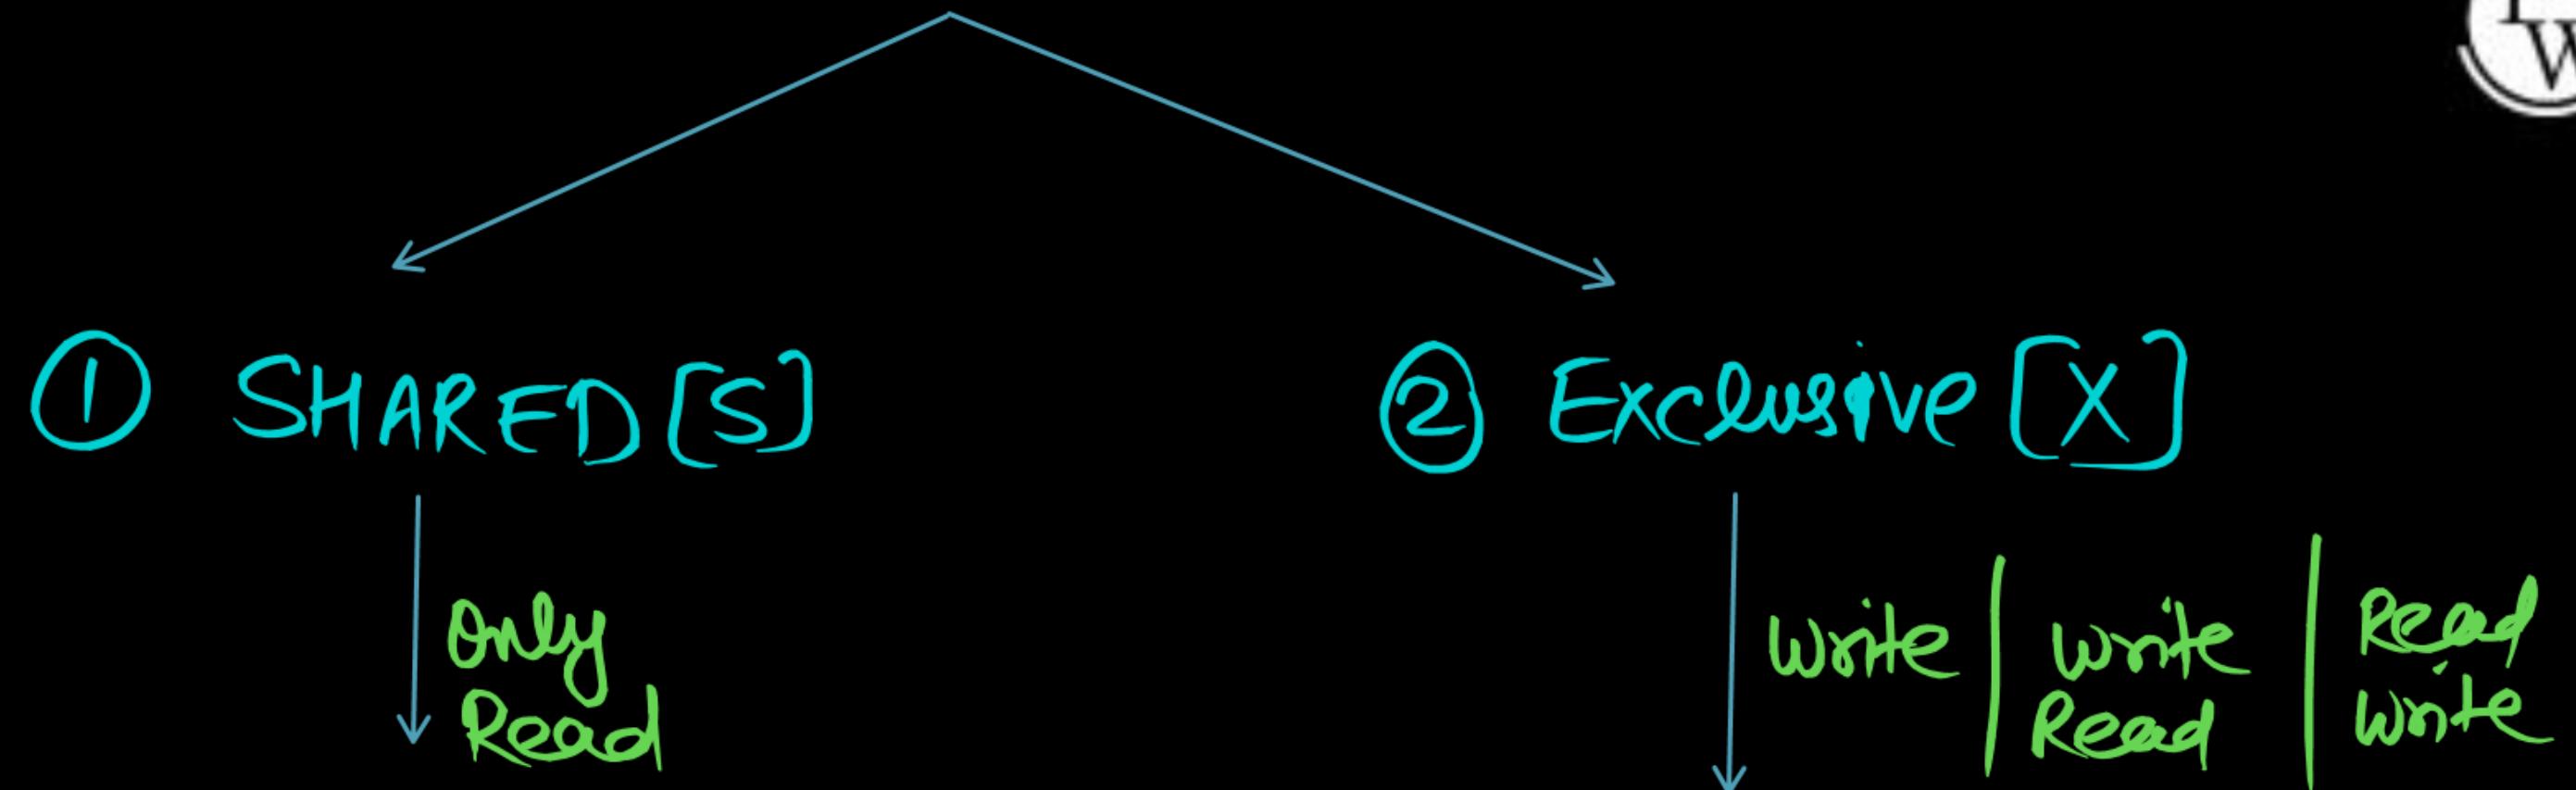




Implementation of Concurrency Control Scheme

LOCK BASED Protocol:





Syntax

LOCK -S(A)
Read (A)
Unlock -S(A)

Syntax

LOCK -X(A)
write (A) | write(A) | Read(A)
Read(A) | write(A)
Unlock -X (A)

P
W

R: Read
W: Write

S: → Only Read
X → Write

		<u>Tj</u>	
		<u>A)</u>	X
<u>Ti</u>	S	<u>S</u>	YES
	X	<u>NO</u>	NO

<u>Ti</u>	<u>Tj</u>	
S(A)	S(A)	✓
S(A)	X(A)	✗
X(A)	S(A)	✗
X(A)	X(A)	✗

Lock-Based Protocols

- ❑ A lock is a mechanism to control concurrent access to a data item
- ❑ Data items can be locked in two modes:
 1. exclusive (X) mode. Data item can be both reads as well as written. X-lock is requested using **lock-X** instruction.
 2. Shared (S) mode. Data item can only be read. S-lock is requested using **lock-S** instruction.
- ❑ Lock requests are made to concurrency-control manager. Transaction can proceed only after request is granted.

Lock-Based Protocols (Cont.)

□ Lock-compatibility matrix

T_j

T_i

DATA ITEM (A)		S	X
Ti	S	true ✓	false ✗
X	false ✗	false ✗	

- A transaction may be granted a lock on an item if the requested lock is compatible with lock already held on the item by other transactions
- Any number of transactions can hold shared locks on an item
- But if any transaction holds an exclusive lock on the item no other transaction may hold any lock on the item.

<u>T₁</u>	<u>T₂</u>	<u>T₁</u>	<u>T₂</u>	<u>LOCK manager</u>
Read(A)		LOCK-X(A)		GRANT-X(A, T _L)
A=A-100		Read(A)		Release /
write(A)		A=A-100		Revoke-X(A, T _L)
	Read(A)	Write(A)		
	Read(C)	Unlock-X(A)		
Read(B)		LOCK-S(A)		GRANT-S(A, T ₂)
B=B+100		Read(A)		Revoke-S(A, T ₂)
write(B)		Unlock-S(A)		
		LOCK-S(C)		GRANT-S(C, T ₂)
		Read(C)		Revoke-S(C, T ₂)
		Unlock-S(C)		
		LOCK-X(B)		GRANT-X(B, T _L)
		R(B)		Revoke-X(B, T _L)
		B=B+100		
		Write(B)		
		Unlock-X(B)		

Schedule with Lock Grants

- A **locking protocol** is a set of rules followed by all transactions while requesting and releasing locks.
- Locking protocols enforce serializability by restricting the set of possible schedules.

T ₁	T ₂	Concurrency-control manager
lock-X(B)		grant-X(B, T ₁)
read(B)		
B:=B - 50		
write(B)		
unlock(B)		
	lock-S(A, T ₂)	revoke-X(B, T ₁)
	read(A)	grant-S(A, T ₂)
	unlock(A)	
	lock-S(B)	revoke-S(A, T ₂)
	read(B)	grant-S(B, T ₂)
	unlock(B)	
	display(A+B)	revoke-S(B, T ₂)
	lock-X(A)	grant-X(A, T ₁)
	read(A)	
	A:=A + 50	
	write(A)	
	unlock(A)	revoke-X(A, T ₁)

Two Phase locking Protocol (2PL)

Lock & Unlock Request done in the 2 Phase.



①

Growing Phase [Acquire lock]



②

Shrinking Phase [Release lock]



Each Transaction first Finish its Growing Phase
then Shrinking Phase.

Two Phase locking Protocol (2PL)

- ① Growing Phase : In the Growing Phase, Transaction May obtain the locks, But Must Not Release Any lock .
- ② Shrinking Phase : In the Shrinking Transaction May Release the lock BUT Must Not ask for Any New lock Request.

T_1	T_2
$\text{Read}(A)$ $A = A - 100$ $\text{Write}(A)$	
<u>$\text{Read}(A)$</u> $\text{Read}(C)$	L.P 
$\text{Read}(B)$ $B = B + 100$ $\text{Write}(B)$	

T_1	T_2	<u>L.M</u>
$\text{LOCK}-X(A)$		$\text{GRANT}-X(A, T_1)$
$\text{Read}(A)$ $A = A - 100$ $\text{Write}(A)$		$\times \text{NO(WAIT)}$
	$\text{Lock} - S(A)$	$\text{GRANT}-X(B, T_1)$
	$\text{Unlock}-X(A)$	$\text{Revoke}-X(A, T_1)$
	$\text{Lock} - S(A)$	$\text{GRANT}-S(A, T_2)$
	$\text{Read}(A)$	
	$\text{Lock} - S(C)$	$\text{GRANT}-S(C, T_2)$
	$\text{Unlock}-S(A)$	$\text{Revoke}-S(A, T_2)$
	$\text{Read}(C)$	$\text{Revoke}-S(C, T_2)$
$\text{Read}(B)$ $B = B + 100$ $\text{Write}(B)$	$\text{Unlock}-S(C)$	
	$\text{Unlock}-X(B)$	$\text{Revoke}-X(B, T_1)$

Important Point about 2PL.

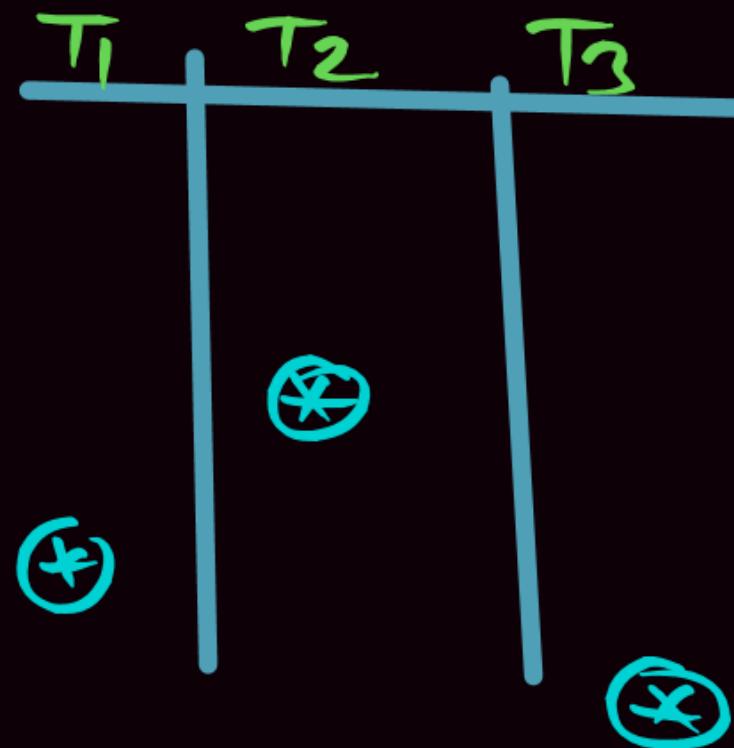
- ① 2PL Ensure Conflict Serializable Schedule
 ↳ Serializability.
- ② If a Schedule is followed by 2PL then Ensure
 Conflict Serializable Schedule.
- ③ Equivalent Serial Schedule [Serializability order] is
 determined by LOCK POINT.

Important Point about PL.

④ LOCK POINT : Position of LAST Lock **or**
First Unlock of the Transaction.

OR

from where Shrinking Phase of the transaction Start.



* First
Unlock

Serializability:
order $\langle T_2, T_1, T_3 \rangle$.

..

Important Point about 2PL.

⑤ 2PL Not Ensure Recoverability.

(Irrecoverable schedule followed by 2PL)

T_1	T_2
$R(A)$ <u>$W(A)$</u>	
	$X(A)$
	$R(A)$
	$W(A)$
	Unlock - $X(A)$
<u>$R(A)$</u>	
Commit	
Commit	
Irrecoverable Schedule.	

Irrecoverable Schedule
but followed by 2PL.

Important Point about PL.

①

T_1	T_2
$R(A)$	
	$w(B)$
$R(B)$	
	<u>$w(A)$</u>

Denied
by T_2

$S(B)$

T_1	T_2
$S(A)$	
$R(A)$	
	$X(B)$
	$w(B)$

Denied
by T_1

$X(A)$

Deadlock

②

T_1	T_2
$w(A)$	
	$R(B)$
$w(B)$	
	$R(A)$

Denied
by T_2

$X(B)$

Deadlock

T_1	T_2
$X(A)$	
$w(A)$	

$S(B)$

$R(B)$

$S(A)$

Denied
by T_1

- ⑥ 2PL Suffers from Deadlock (Not Free From Deadlock)
- ⑦ 2PL Suffers from Starvation.

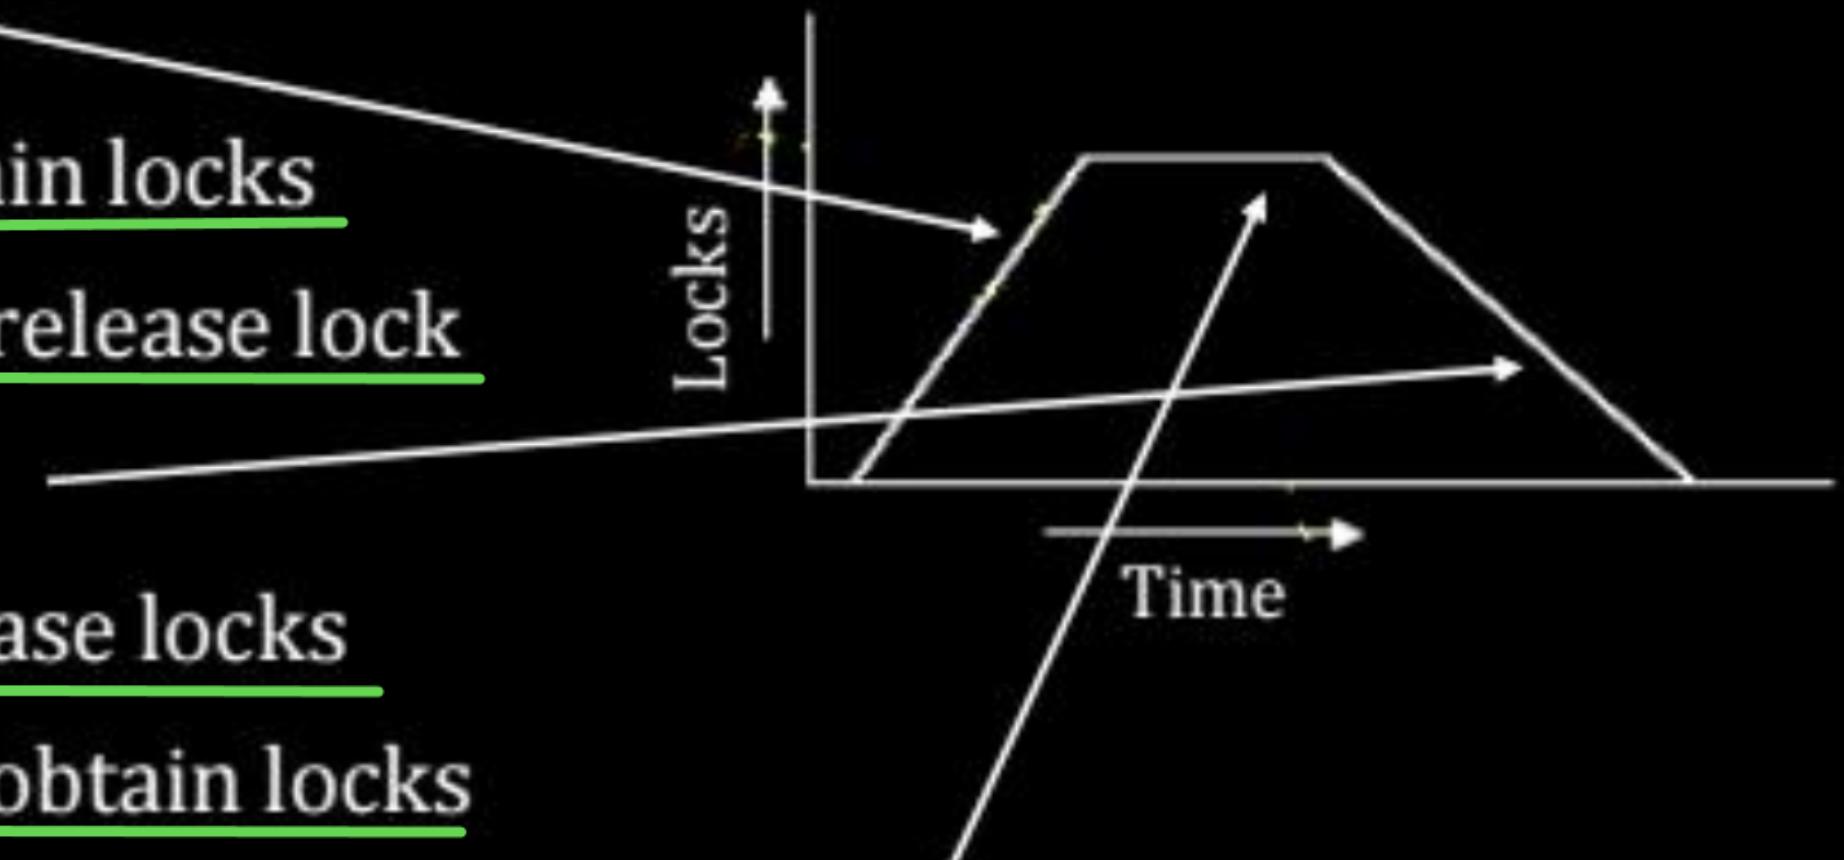
T_1	T_2	T_3	T_4
	<u>$S(A)$</u>		
Deny by $T_2 \rightarrow X(A)$		$S(A)$	
Deny by $T_3 \rightarrow X(A)$	$Unlock(A)$		
Deny by $T_4 \rightarrow X(A)$		$unlock_s(A)$	$S(A)$
granted $X(A)$			$unlock_s(A)$

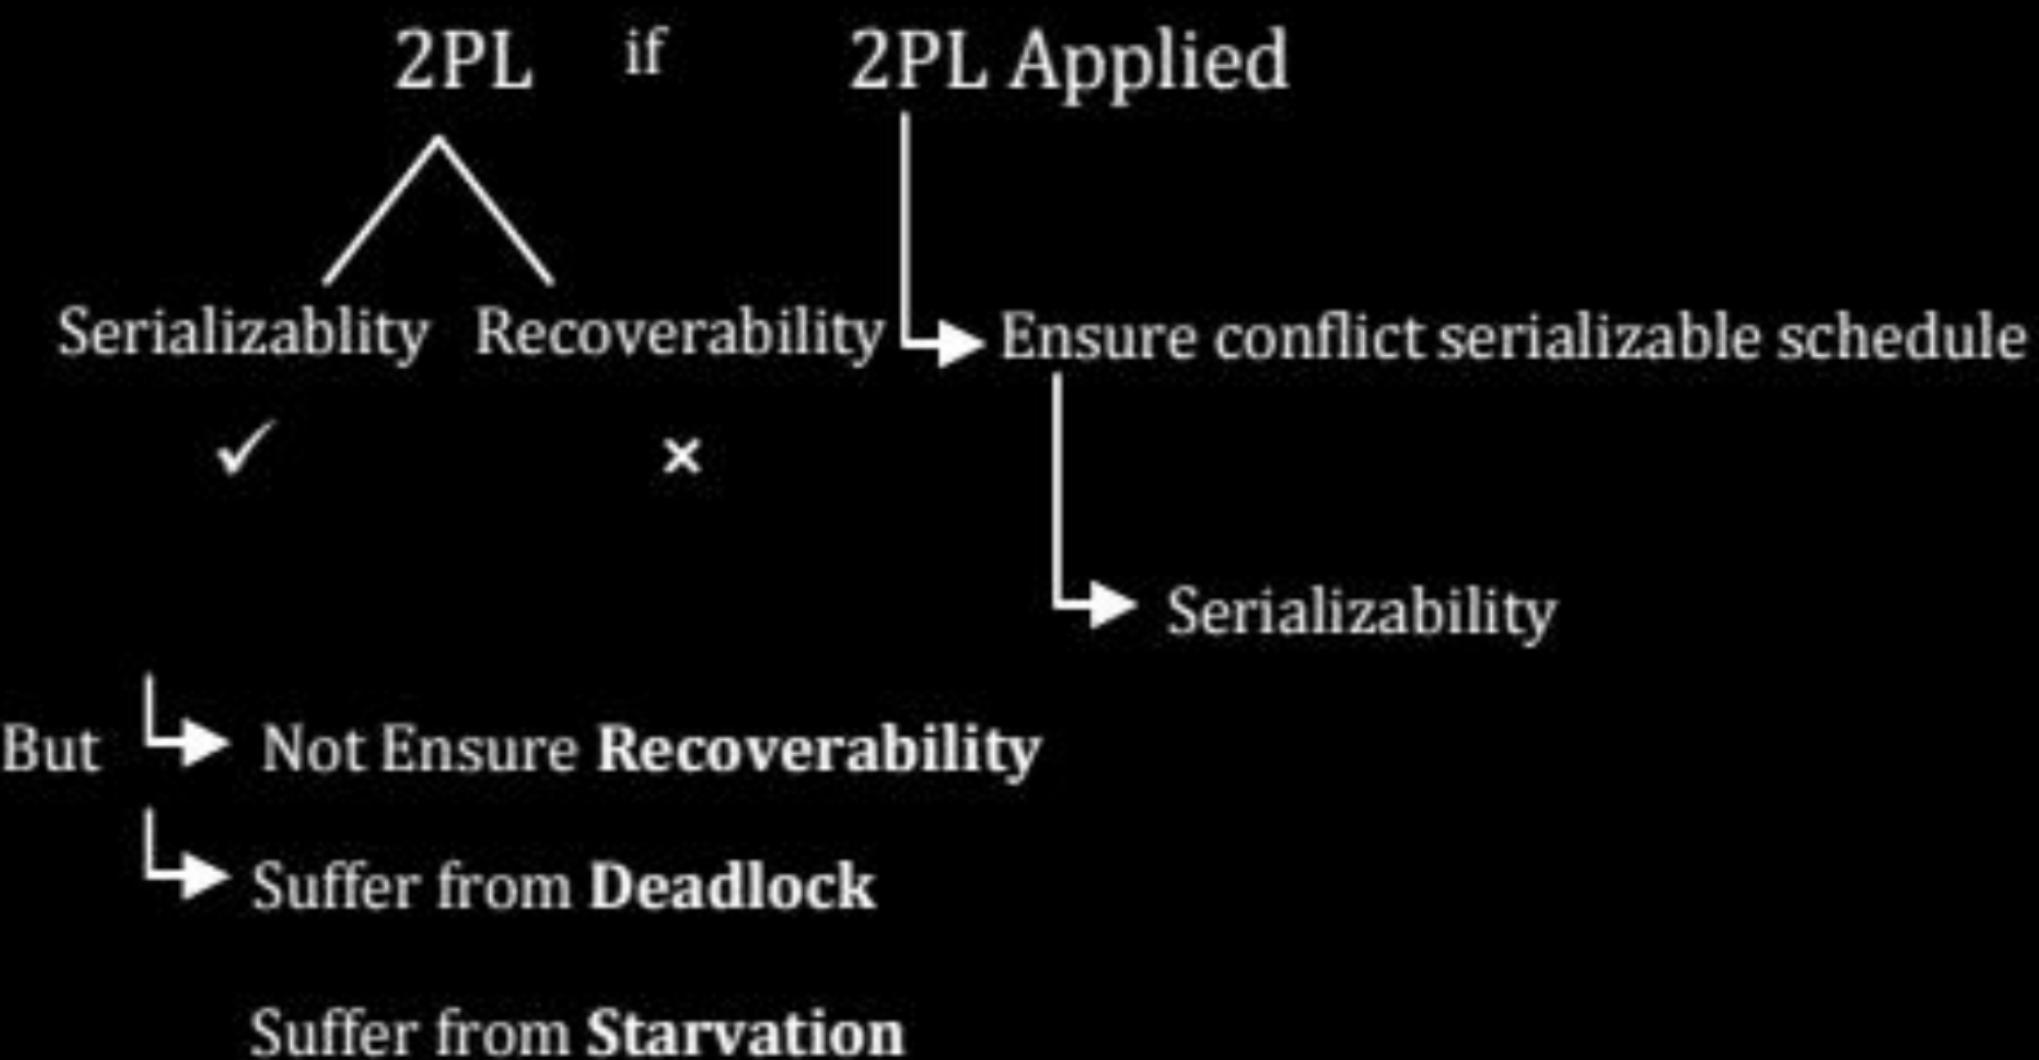
2PL Suffer
from Starvation.

...

The Two-Phase Locking Protocol

- A protocol which ensures conflict-serializable schedules.
- Phase 1: Growing Phase
 - ❖ Transaction may obtain locks
 - ❖ Transaction may not release lock
- Phase 2: Shrinking Phase
 - ❖ Transaction may release locks
 - ❖ Transaction may not obtain locks
- The protocol assures serializability. It can be proved that transactions can be serialized in the order of their lock points (i.e., the point where a transaction acquired its final lock).





STRICT 2PL.:

2PL + All Exclusive Lock taken by transaction Must be Release After Commit **or** Rollback .

T ₁	T ₂
X(A)	
w(A)	
C R	
unlock-X(A)	S(A) / X(A) R(A) / w(A)

Strict Recoverable.

OR

2PL + All Exclusive lock taken by the Transaction Must be Held untill Commit | Rollback of the Transaction.

T _i
-X(A)
S(B)
-X(C)
S(D)
≡
Unlock -(R)
Unlock -(D)
C R
-unlock-X(A)
-unlock-X(C)

T_1	T_2
w(A)	
CLR	R(A)

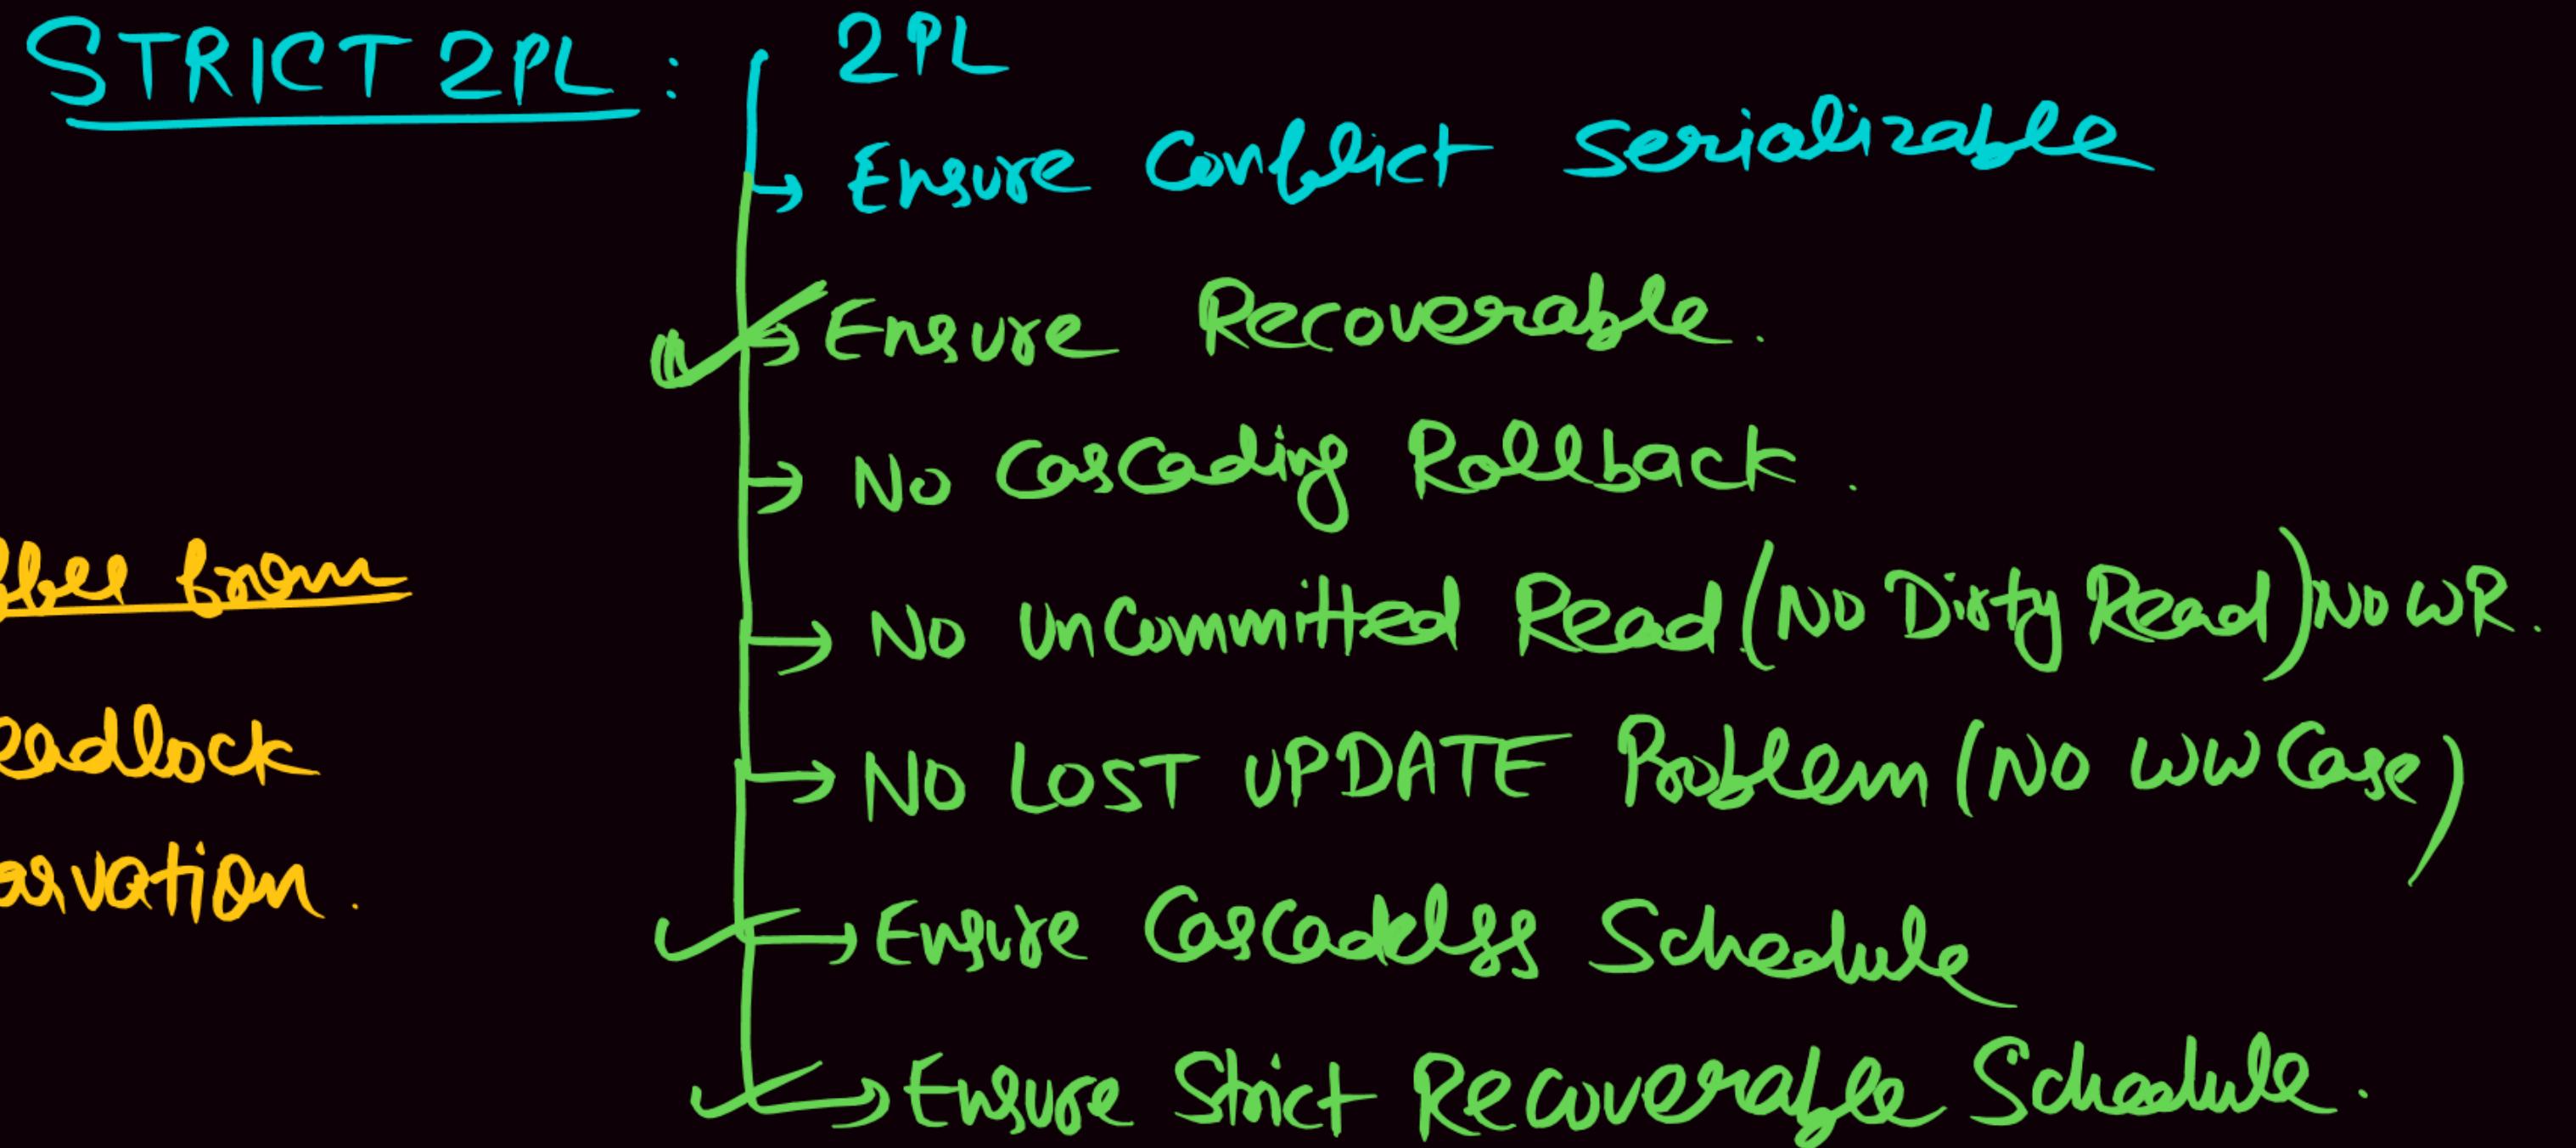
Recoverable

T_1	T_2
w(A)	
CLR	R(A)

Consistent

T_1	T_2
w(A)	
CLR	S(A) / X(A) R(A) / w(A)

Strict Recoverable



Subbe from

- Deadlock
- Starvation .

Rigorous 2PL

: 2PL + ALL Locks (SHARED & Exclusive)



taken by the transaction Must be Held
by the transaction Until Commit / Rollback .

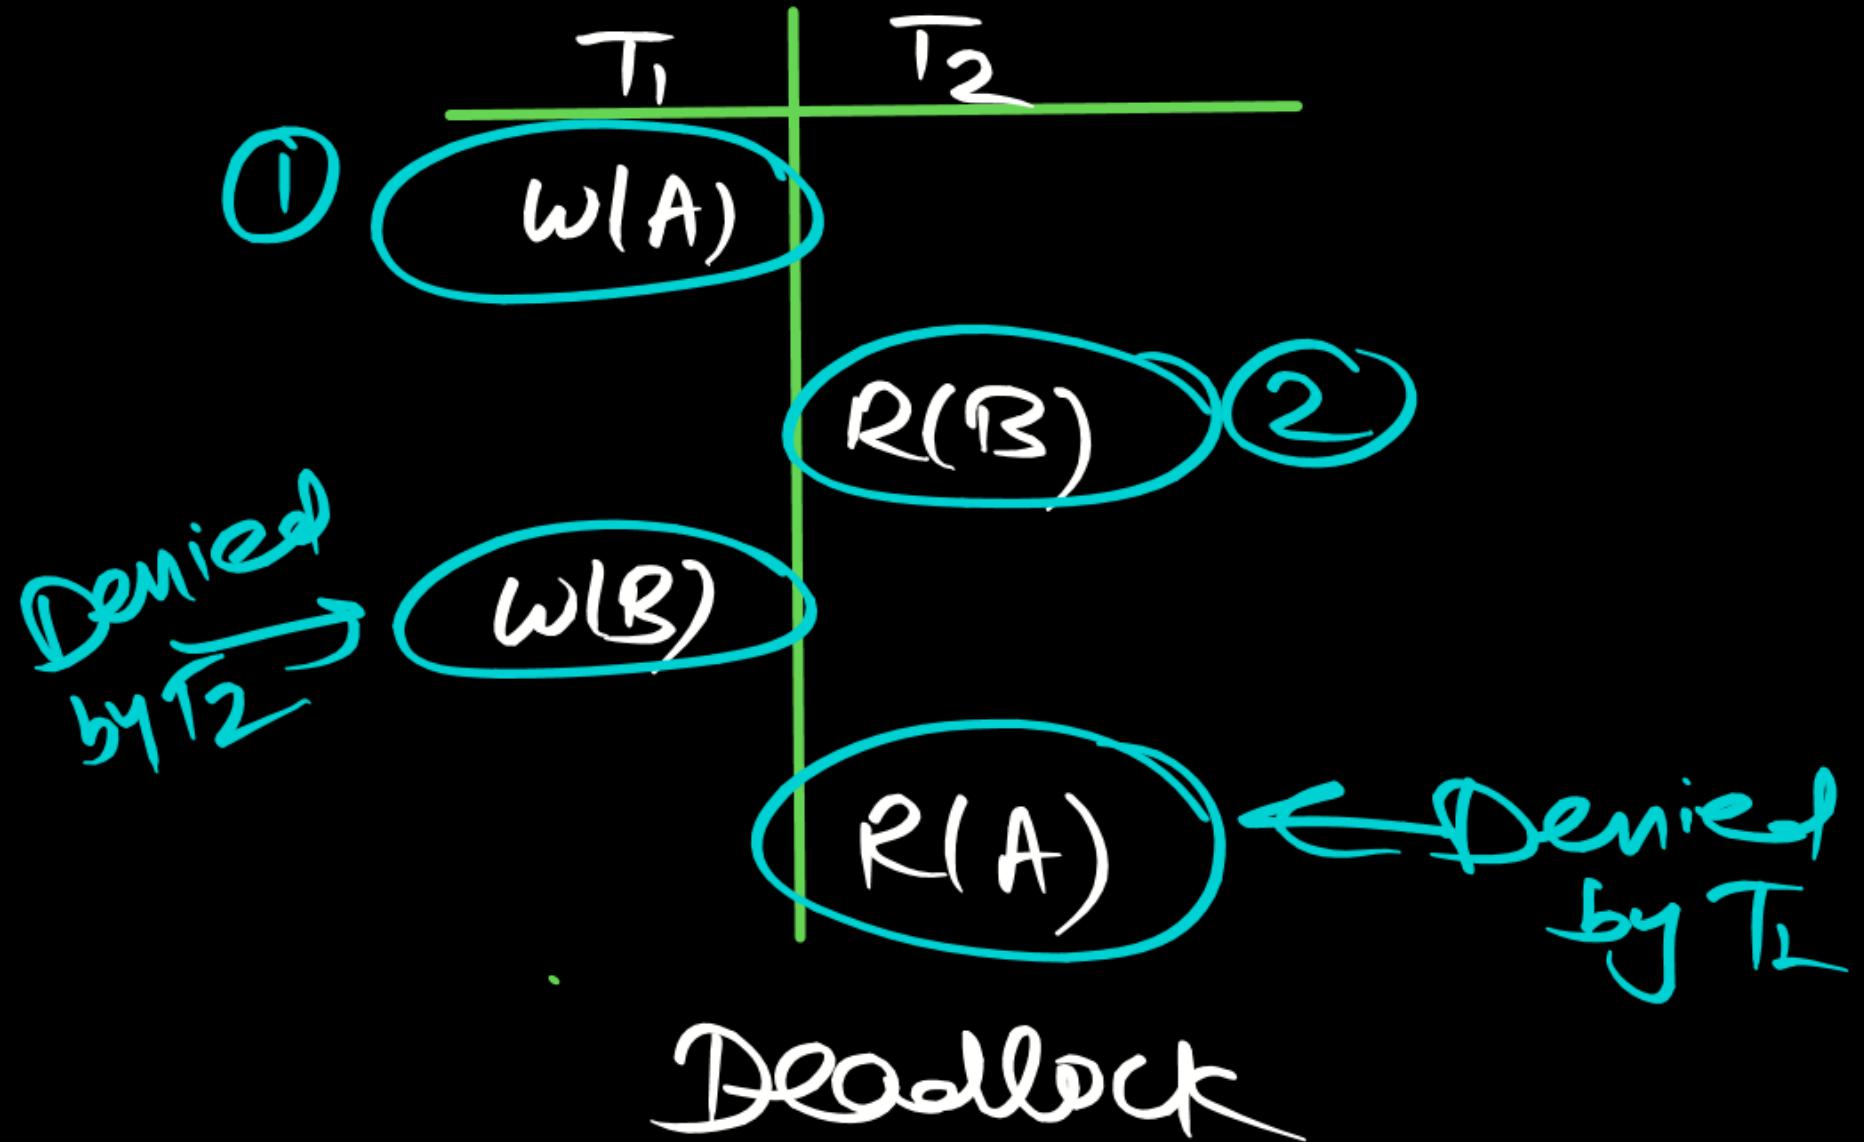
OR

Release After Commit | Rollback .

<u>T_i</u>
X(A)
S(B)
X(C)
S(D)
≡
C R
U(A)
U(B)
U(C)
U(D)

Rigorous 2PL Subber from .

- ① Deadlock
- ② Starvation .



Conservative 2PL : Each Transaction acquire the Lock in the beginning of the Transaction (Before Start) & Release all lock After Commit / Rollback .

Conservative 2PL → Ensure Conflict Serializable
→ Recoverable, Cascaderless, Strict Recoverable
→ No Cascading Rollback / Abort
→ No Uncommit Read | ^{No} Dirty Read
→ No Lost UPDATE Problem
↳ NO DEADLOCK .

Subbed from
STARVATION

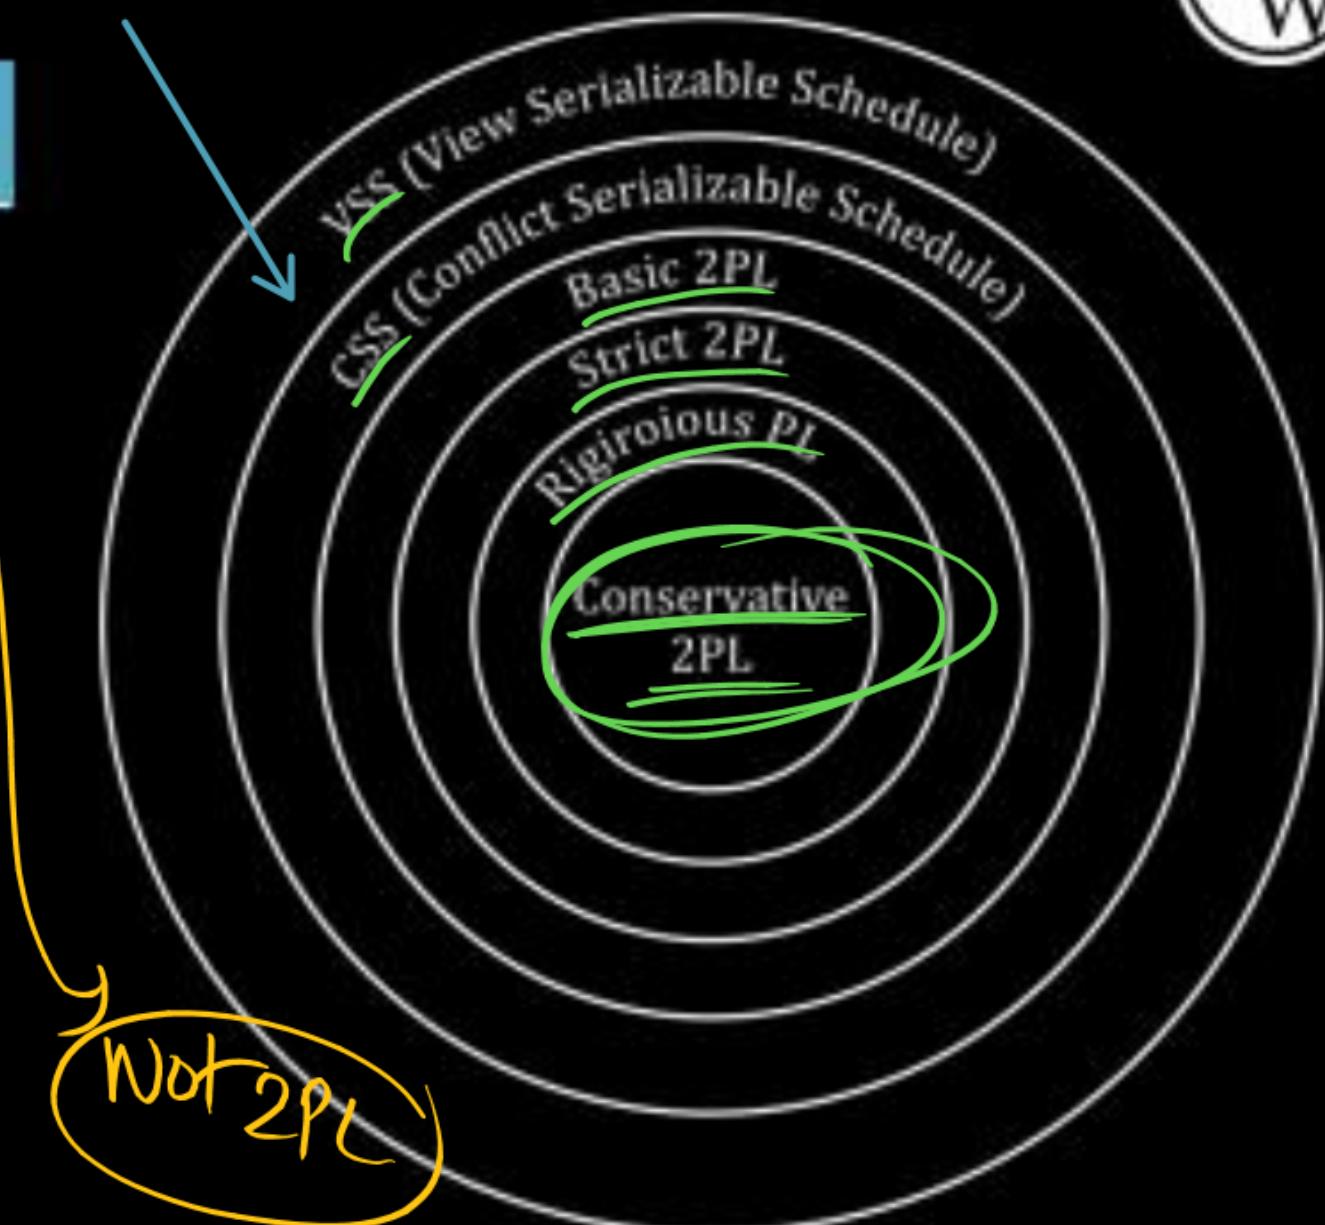
The Two-Phase Locking Protocol (Cont.)

- Two-phase locking does not ensure freedom from deadlocks
- Extensions to basic two-phase locking needed to ensure recoverability of freedom from cascading roll-back
 - * ② Strict two-phase locking: a transaction must hold all its exclusive locks till it commits/aborts.
 - Ensures recoverability and avoids cascading roll-backs
 - * ③ Rigorous two-phase locking: a transaction must hold all locks till commit/abort.
 - Transactions can be serialized in the order in which they commit.
- Most databases implement rigorous two-phase locking, but refer to it as simply two-phase locking

④ Conservative
2PL

1	2	3	4	5
Lock-S(A) R(A) <u>Lock-X(B)</u> R(B) <u>Unlock(A)</u> W(B) <u>Unlock(B)</u> Commit	Lock-S(A) R(A) <u>Lock-X(B)</u> <u>Unlock(A)</u> R(B) W(B) Commit	Lock-S(A) R(A) <u>Lock-X(B)</u> R(A) R(B) W(B) <u>Commit</u>	Lock-S(A) R(A) Lock-X(B) R(A) R(B) W(B) <u>Commit</u> <u>Unlock(A)</u> Unlock(B)	Lock-S(A) R(A) Unlock(A) Lock-X(B) R(B) R(B) W(B) R(B) W(B) Unlock(B) Unlock(B)
BASIC 2PL	Strict 2PL (Basic 2PL)	Rigorous 2PL (Strict 2PL Basic 2PL)	Congervative 2PL (Rigorous 2PL Strict 2PL Basic 2PL)	

Congervative



TIMESTAMP BASED CONCURRENCY CONTROL

Timestamp-Based Protocols

- ❑ Each transaction T_i is issued a timestamp $TS(T_i)$ when it enters the system.
 - ❖ Each transaction has a unique timestamp
 - ❖ Newer transaction have timestamp strictly greater than earlier ones
 - ❖ Timestamp could be based on a logical counter
 - Real time may not be unique
 - Can use (wall-clock time, logical counter) to ensure
- ❑ Timestamp-based protocols manage concurrent execution such that **time-stamp order = serializability order**
- ❑ Several alternative protocols based on timestamps

Timestamp-Based Protocols

The **timestamp ordering (TSQ) protocol**

- Maintains for each data Q two timestamp values:
 - ❖ **W-timestamp(Q)** is the largest time-stamp of any transaction that executed **write** (Q) successfully.
 - ❖ **R-timestamp (Q)** is the largest time-stamp of any transaction that executed **read** (Q) successfully.
- Imposes rules on read and write operations to ensure that
 - ❖ any conflicting operations are executed in timestamp order
 - ❖ out of order operations cause transaction rollback

I: T_i - Read(Q) (Transaction T_i Issue R(Q) Operation)

- (i) If $TS(T_i) < WTS(Q)$: Read operation Reject & T_i Rollback.
- (ii) If $TS(T_i) \geq WTS(Q)$: Read operation is allowed
and Set $Read - TS(Q) = \max[RTS(Q), TS(T_i)]$

II: T_i - Write(Q) (Transaction T_i Issue Write(Q) Operation)

- (i) If $TS(T_i) < RTS(Q)$: Write operation Reject & T_i Rollback.
- (ii) If $TS(T_i) < WTS(Q)$: Write operation Reject & T_i Rollback.
- (iii) Otherwise execute write (Q) operation
Set $Read WTS(Q) = TS(T_i)$

Timestamp-Based Protocols (Cont.)

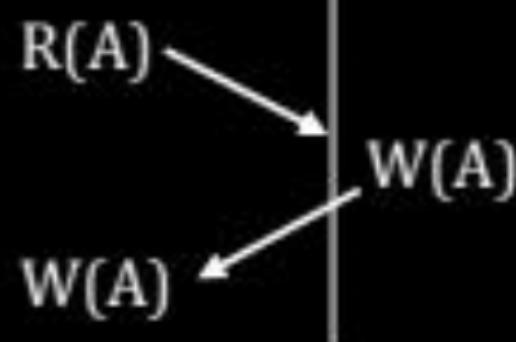
- Suppose a transaction T_i issues a **read** (Q)
 1. If $TS(T_i) \leq W\text{-timestamp } (Q)$, then T_i needs to read a value of Q that was already overwritten.
 - Hence, the **read** operation is rejected, and T_i is rolled back.
 2. If $TS(T_i) \geq W\text{-timestamp } (Q)$, then the **read** operation is executed, and $R\text{-timestamp}(Q)$ is set to.
 $\max(R\text{-timestamp } (Q), TS (T_i))$.

Timestamp-Based Protocols (Cont.)

- Suppose a transaction T_i issues **write(Q)**
 1. If $TS(T_i) < R\text{-timestamp}(Q)$, then the value of Q that T_i is producing was needed previously, and the system assumed that the value would never be produced.
 - Hence, the **write** operation is rejected, and T_i is rolled back.
 2. If $TS(T_i) < W\text{-timestamp}(Q)$, then T_i is attempting to write an obsolete value of Q.
 - Hence, this **write** operation is rejected, and T_i is rolled back.
 3. Otherwise, the **write** operation is executed, and $W\text{-timestamp}(Q)$ is set to $TS(T_i)$.

(1)

$T_1(10)$	$T_2(20)$
-----------	-----------

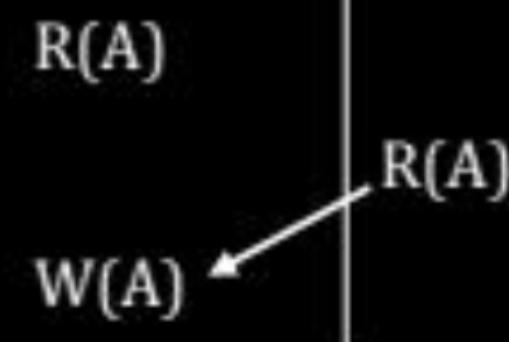


$TS(T_1) < TS(T_2)$

$T_1 \rightarrow T_2$

(2)

$T_1(10)$	$T_2(20)$
-----------	-----------



$TS(T_1) < TS(T_2)$

$T_1 \rightarrow T_2$

(3)

$T_1(10)$	$T_2(20)$
-----------	-----------

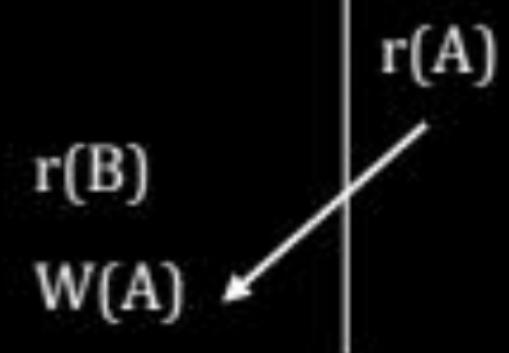


$TS(T_1) < TS(T_2)$

$T_1 \rightarrow T_2$

(4)

$T_1(10)$	$T_2(20)$
-----------	-----------



$TS(T_1) < TS(T_2)$

$T_1 \rightarrow T_2$

Thomas' Write Rule

- ❑ Modified version of the timestamp-ordering protocol in which obsolete **write** operations may be ignored under certain circumstances.
- ❑ When T_i attempts to write data item Q , if $TS(T_i) < W\text{-timestamp}(Q)$, then T_i is attempting to write an obsolete value of $\{Q\}$.
 - ❖ Rather than rolling back T_i as the timestamp ordering protocol would have done, this **{write}** operation can be ignored.
- ❑ Otherwise this protocol is the same as the timestamp ordering protocol.
- ❑ Thomas' Write Rule allows greater potential concurrency.
 - ❖ Allows some view-serializable schedules that are not conflict-serializable.

Thomas Write Rule (View Serializability)

1. $TS(T_i) < RTS(Q)$: Rollback
2. $TS(T_i) < WTS(Q)$: Write operation is Ignored and No Roll back

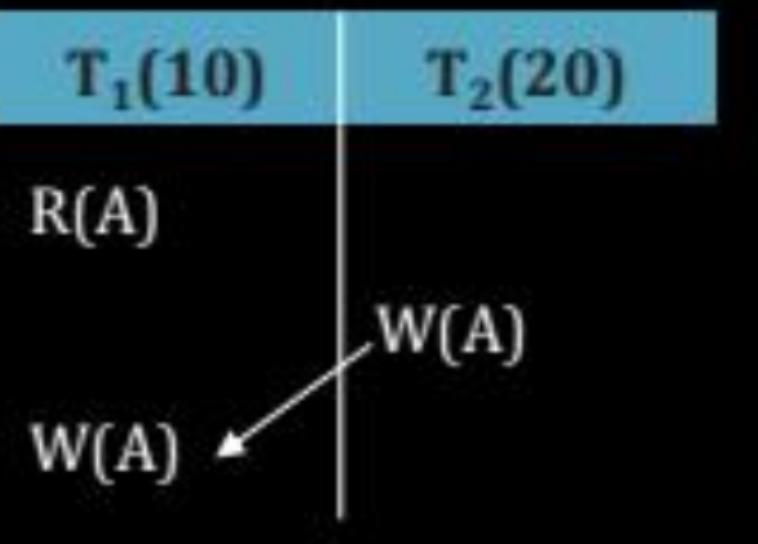
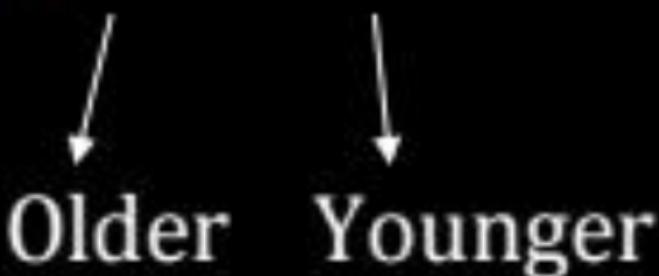
Same as TSP

Time Stamp Protocol: Ensure serializability deadlock free but starvation possible

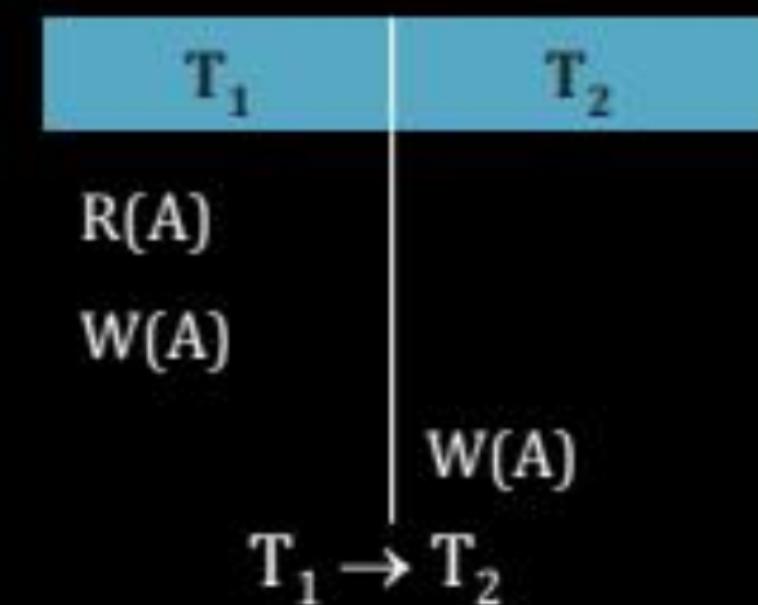
Deadlock Prevention Algorithm

(1) Wait-Die

(2) Wound-wait



$TS(T_1) < TS(T_2)$
 $T_1 \rightarrow T_2$



$T_1 \rightarrow T_2$

Q.

In which one of the following Lock Scheme Deadlock cannot occur?

[MCQ]

P
W

- A** Basic 2PL
- B** Strict 2PL
- C** Conservative 2PL
- D** Rigorous 2PL

Q.

Consider the following statement about lock-based protocol

P
W

- (A) 2 PL (2phase locking) protocol Ensure view serializability
- (B) 2PL ensure recoverability & No cascading rollback.
- (C) Strict 2 PL ensure recoverability & no cascading rollback.
- (D) Strict 2 PL avoids deadlock (not suffering from deadlock).

How many numbers of above statement are correct?

A

1

[MCQ]

B

2

C

3

D

4

Q.

Consider the following Schedule:

$r_1(x) \ r_2(y) \ r_2(x) \ w_1(z) \ r_1(y) \ w_3(y) \ r_3(z) \ w_2(y) \ w_3(x)$

which of the following time stamp ordering Not allows to execute schedule using Thomas Write rule time stamp Ordering Protocol?

[MSQ]

P
W

- A** $(T_1, T_2, T_3) = (20, 30, 10)$
- B** $(T_1, T_2, T_3) = (10, 20, 30)$
- C** $(T_1, T_2, T_3) = (10, 30, 20)$
- D** $(T_1, T_2, T_3) = (30, 20, 10)$

Q.

Consider the following statements:

S₁: All strict recoverable schedules are serial.

S₂: All recoverable schedules are conflict serializable.

S₃: All strict schedules are conflict serializable.

S₄: All conflict serializable schedules are free from cascading rollbacks.

Which of the following is true?

- (a) Only S₁ and S₄
- (b) Only S₂, S₃ and S₄
- (c) Only S₂ and S₄
- (d) None of these

Q.

Consider the following transaction:

$T_1: R_1(x) W_1(x) R_1(y) W_1(y)$

$T_2: W_2(y) W_2(x)$

The number of non-serial schedules between T_1 and T_2 which are serializable?

(a) 2

(b) 13

(c) 15

(d) None of these

Any Doubt ?

**THANK
YOU!**

